

# OpenMP case studies

Victor Eijkhout





A particle has  $x, y$  coordinates and a mass  $c$ . For two particles  $(x_1, y_1, c_1)$ ,  $(x_2, y_2, c_2)$  the force on particle 1 from particle 2 is:

$$\vec{F}_{12} = \frac{c_1 \cdot c_2}{\sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}} \cdot \vec{r}_{12}$$

where  $\vec{r}_{12}$  is the unit vector pointing from particle 2 to 1. With  $n$  particles, each particle  $i$  feels a force

$$\vec{F}_i = \sum_{j \neq i} \vec{F}_{ij}.$$



Let's start with a couple of building blocks.

```
1 // molecularstruct.c
2 struct point{ double x,y; double c; };
3 struct force{ double x,y; double f; };
4
5 /* Force on p1 from p2 */
6 struct force force_calc( struct point p1,struct point p2 ) {
7     double dx = p2.x - p1.x, dy = p2.y - p1.y;
8     double f = p1.c * p2.c / sqrt( dx*dx + dy*dy );
9     struct force exert = {dx,dy,f};
10    return exert;
11 }
```



(Probably wrong, but hey, I'm not a physicist)

```
1 void add_force( struct force *f,struct force g ) {  
2     f->x += g.x; f->y += g.y; f->f += g.f;  
3 }  
4 void sub_force( struct force *f,struct force g ) {  
5     f->x -= g.x; f->y -= g.y; f->f += g.f;  
6 }
```

For reference, this is the sequential code:

```
1 for (int ip=0; ip<N; ip++) {  
2     for (int jp=ip+1; jp<N; jp++) {  
3         struct force f = force_calc(points[ip],points[jp]);  
4         add_force( forces+ip,f );  
5         sub_force( forces+jp,f );  
6     }  
7 }
```

Here  $\vec{F}_{ij}$  is only computed for  $j > i$ , and then added to both  $\vec{F}_i$  and  $\vec{F}_j$ .



In C++ we can have a class with an addition operator and such:

```
1 // molecular.cxx
2 class force {
3 private:
4     double _x{0.},_y{0.}; double _f{0.};
5 public:
6     force() {};
```

```
7     force(double x,double y,double f)
8         : _x(x),_y(y),_f(f) {};
```

```
1 force operator+( const force& g ) {
2     return { _x+g._x, _y+g._y, _f+g._f };
3 }
```

Sequential code:

```
1 for (int ip=0; ip<N; ip++) {
2     for (int jp=ip+1; jp<N; jp++) {
3         force f = points[ip].force_calc(points[jp]);
4         forces[ip] += f;
5         forces[jp] -= f;
6     }
```



Is the outer loop parallelizable? The inner? Both together?



One solution would be to compute the  $\vec{F}_{ij}$  interactions for all  $i, j$ , so that there are no conflicting writes.

```
1  for (int ip=0; ip<N; ip++) {  
2      struct force sumforce;  
3      sumforce.x=0.; sumforce.y=0.; sumforce.f=0.;  
4  #pragma omp parallel for reduction(+:sumforce)  
5      for (int jp=0; jp<N; jp++) {  
6          if (ip==jp) continue;  
7          struct force f = force_calc(points[ip],points[jp]);  
8          sumforce.x += f.x; sumforce.y += f.y; sumforce.f += f.f;  
9      } // end parallel jp loop  
10     add_force( forces+ip, sumforce );  
11 } // end ip loop
```



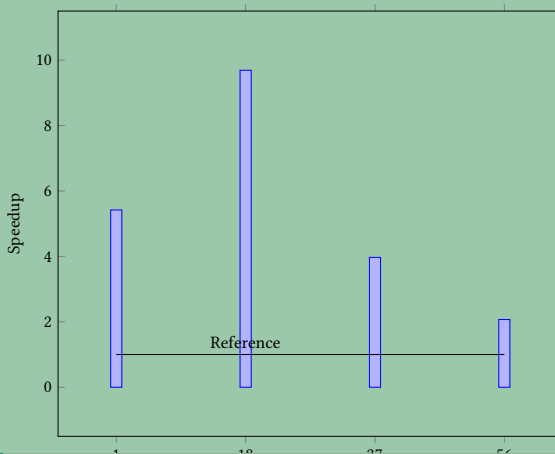


In C++ we use the fact that we can reduce on any class that has an addition operator:

```
1  for (int ip=0; ip<N; ip++) {  
2      force sumforce;  
3      #pragma omp parallel for reduction(+:sumforce)  
4      for (int jp=0; jp<N; jp++) {  
5          if (ip==jp) continue;  
6          force f = points[ip].force_calc(points[jp]);  
7          sumforce += f;  
8      } // end parallel jp loop  
9      forces[ip] += sumforce;  
10 } // end ip loop
```



This increases the scalar work by a factor of two, but surprisingly, on a single thread the run time improves: we measure a speedup of 6.51 over the supposedly 'optimal' code. (Why?)




The *i* update is fine, we make the *j* update atomic:

```
1  #pragma omp parallel for schedule(guided,4)
2      for (int ip=0; ip<N; ip++) {
3          for (int jp=ip+1; jp<N; jp++) {
4              struct force f = force_calc(points[ip],points[jp]);
5              add_force( forces+ip,f );
6              sub_force( forces+jp,f );
7          }
8      }
```

To deal with the conflicting *jp* writes, we make the writes atomic:

```
1  void sub_force( struct force *f,struct force g ) {
2      #pragma omp atomic
3          f->x -= g.x;
4      #pragma omp atomic
5          f->y -= g.y;
6      #pragma omp atomic
7          f->f += g.f;
8      }
```






What happens with one thread?

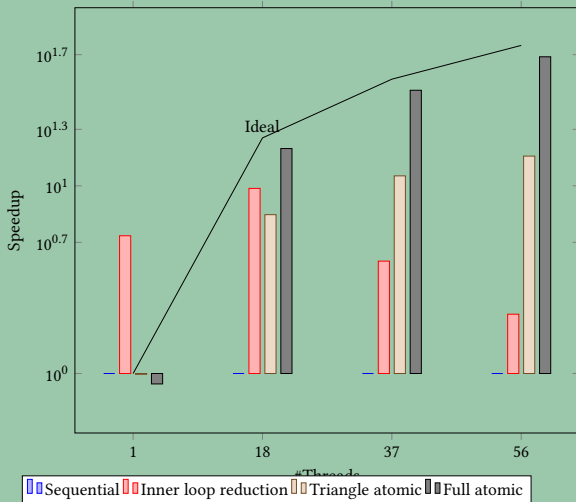


But if we decide to use atomic updates, we can take the full square loop, collapse the two loops, and make every write atomic.

```
1  #pragma omp parallel for collapse(2)
2      for (int ip=0; ip<N; ip++) {
3          for (int jp=0; jp<N; jp++) {
4              if (ip==jp) continue;
5              struct force f = force_calc(points[ip],points[jp]);
6              add_force( forces+ip, f );
7          } // end parallel jp loop
8      } // end ip loop
```












Search: traverse the tree,  
and abort unsuccessful branches

DFS, not BFS



```
1 placement initial; initial.fill(empty);
2 auto solution = place_queen(0,initial);
3
4 optional<placement> place_queen
5     (int iqueen,const placement& current) {
6     for (int col=0; col<N; col++) {
7         placement next = current;
8         next.at(iqueen) = col;
9         if (feasible(next)) {
10             if (iqueen==N-1)
11                 return next;
12             auto attempt = place_queen(iqueen+1,next);
13             if (attempt.has_value())
14                 return attempt;
15         } // end if(feasible)
16     }
17     return {};
18 };
```



```
1 placement initial; initial.fill(empty);  
2 optional<placement> eightqueens;  
3 #pragma omp parallel  
4 #pragma omp single  
5 eightqueens = place_queen(0,initial);
```



We create a task for each column, and since they are in a loop we use `taskgroup` rather than `taskwait`.

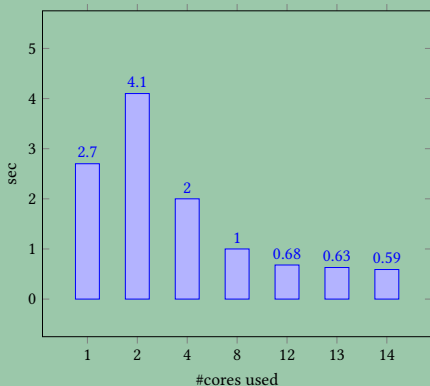
```
1  #pragma omp taskgroup
2    for (int col=0; col<N; col++) {
3        placement next = current;
4        next.at(iqueen) = col;
5  #pragma omp task firstprivate(next)
6        if (feasible(next)) {
7            // stuff
8        } // end if(feasible)
9    }
```



However, the sequential program had `return` and `break` statements in the loop, which is not allowed in workshare constructs such as `taskgroup`. Therefore we introduce a return variable, declared as shared:

```
1  // queens0.cxx
2  optional<placement> result = {};
3  #pragma omp taskgroup
4  for (int col=0; col<N; col++) {
5      placement next = current;
6      next.at(iqueen) = col;
7      #pragma omp task firstprivate(next) shared(result)
8      if (feasible(next)) {
9          if (iqueen==N-1) {
10             result = next;
11         } else { // do next level
12             auto attempt = place_queen(iqueen+1,next);
13             if (attempt.has_value()) {
14                 result = attempt;
15             }
16         } // end if(iqueen==N-1)
17     } // end if(feasible)
18 }
```






This is a 1000 times slower than sequential. Why?



Body of the loop over columns:

```
1  // queenfinal.cxx
2  if (feasible(next)) {
3      if (iqueen==N-1) {
4          result = next;
5          #pragma omp cancel taskgroup
6      } else { // do next level
7          auto attempt = place_queen(iqueen+1,next);
8          if (attempt.has_value()) {
9              result = attempt;
10             #pragma omp cancel taskgroup
11         }
12     } // end if (iqueen==N-1)
13 } // end if (feasible)
```





Still not great. Conclusion?







Parallel loops in C++ can use range-based syntax:

```
1 // speedup.cxx
2 #pragma omp parallel for
3 for ( auto& v : values ) {
4     for (int jp=0; jp<M; jp++) {
5         double f = sin( v );
6         v = f;
7     }
8 }
```

Tests not reported here show exactly the same speedup as the C code.



Support for *C++ iterators*

```
1  #pragma omp declare reduction (merge : std::vector<int>  
2      : omp_out.insert(omp_out.end(), omp_in.begin(), omp_in.end()))
```



You can reduce with a templated function if you put both the declaration and the reduction in the same templated function:

```
1  template<typename T>
2  T generic_reduction( vector<T> tdata ) {
3  #pragma omp declare reduction                                     \
4      (rwzt:T:omp_out=reduce_without_zero<T>(omp_out,omp_in))    \
5      initializer(omp_priv=-1.f)
6
7      T tmin = -1;
8  #pragma omp parallel for reduction(rwzt:tmin)
9      for (int id=0; id<tdata.size(); id++)
10         tmin = reduce_without_zero<T>(tmin,tdata[id]);
11     return tmin;
12 };
```

which is then called with specific data:

```
1  auto tmin = generic_reduction<float>(fdata);
```



Reduction can be applied to any class for which the reduction operator is defined as `operator+` or whichever operator the case may be.

```
1 // reductcomplex.cxx
2 class Thing {
3 private:
4     float x;
5 public:
6     Thing() : Thing( 0.f ) {};
7     Thing( float x ) : x(x) {};
8     Thing operator+( const Thing&
9         ↪other ) {
10         return Thing( x + other.x );
11     };
12 };
13 }
```

```
1 vector< Thing >
2     ↪things(500,Thing(1.f) );
3 Thing result(0.f);
4 #pragma omp parallel for
5     ↪reduction( +:result )
6 for ( const auto& t : things )
7     result = result + t;
```

A default constructor is required for the internally used init value; see figure ??.



```
1  // lockobject.cxx
2  class atomic_int {
3  private:
4      omp_lock_t the_lock;
5      int _value{0};
6  public:
7      atomic_int() {
8          omp_init_lock(&the_lock);
9      };
10     atomic_int( const atomic_int& )
11         = delete;
12     atomic_int& operator=( const atomic_int& )
13         = delete;
14     ~atomic_int() {
15         omp_destroy_lock(&the_lock);
16     };
```

Running this:

```
1  atomic_int my_object;
2  vector<std::thread> threads;
3  for (int ithread=0; ithread<NTHREADS; ithread++) {
4      threads.push_back
```



We make a template for uninitialized types:

```
1 // heatalloc.cxx
2 template<typename T>
3 struct uninitialized {
4     uninitialized() {}
5     T val;
6     constexpr operator T() const {return val;};
7     T operator=( const T&& v ) { val = v; return val; };
8 };
```

so that we can create vectors that behave normally:

```
1 vector<uninitialized<double>> x(N),y(N);
2
3 #pragma omp parallel for
4 for (int i=0; i<N; i++)
5     y[i] = x[i] = 0.;
6 x[0] = 0; x[N-1] = 1.;
```

