

Optimizing random walks

Victor Eijkhout, Susan Lindsey

Fall 2023

last formatted: February 6, 2024

1. Malaria

- A mosquito flies in a straight line for some unit time,
- then it turns in a random direction.
- How far does it get in N time intervals?
- Answer: about \sqrt{N} .

2.

Code:

```
1 // rand/walk_vec.cpp
2 float avg_dist{0.f};
3 for ( int x=0; x<experiments; ++x )
4 {
5     Mosquito m(dim);
6     for (int step=0; step<steps;
7         ++step)
8         m.step();
9     avg_dist += m.distance();
10 }
11 avg_dist /= experiments;
```

Output:

```
D=3 after 10000 steps,
    distance= 83.7997
D=3 after 100000 steps,
    distance= 224.372
D=3 after 1000000
    steps, distance=
    922.599
product took: 2776
    milliseconds
```

3.

```
// rand/walk_lib_vec.cpp
class Mosquito {
private:
    vector<float> pos;
public:
    Mosquito( int d )
        : pos( vector<float>(d,0.f) ) { };
};
```

4.

```
// rand/walk_lib_vec.cpp
void step() {
    int d = pos.size();
    auto incr = random_step(d);
    for (int id=0; id<d; ++id)
        pos.at(id) += incr.at(id);
};
```

5.

```
// rand/walk_lib_vec.cpp
vector<float> random_coordinate( int d ) {
    auto v = vector<float>(d);
    for ( auto& e : v )
        e = random_float();
    return v;
};
```

6.

```
// rand/walk_lib_vec.cpp
vector<float> random_step(int d) {
    for (;;) {
        auto step = random_coordinate(d);
        if ( auto l=length(step); l<=1.f ) {
            if ( l==0.f ) {
                /*
                 * Zero lengths can conceivably happen for d==1
                 * but should not for higher d.
                 */
                assert(d==1);
            } else {
                normalize(step,l);
                return step;
            }
        }
    }
};
```

7. exercise

Take the basic code, and make a version based on

```
template<int d>  
class Mosquito { /* ... */
```

How much does this simplify your code? Do you get any performance improvement?

You can base this off the file `walk_vec.cpp` in the repository

8.

So we move the creation of the vectors outside of the computational routines. The random coordinates are now written into an array passed as parameter:

```
// rand/walk_lib_pass.cpp
void random_coordinate( vector<float>& v ) {
    for ( auto& e : v )
        e = random_float();
};
```

9.

Likewise the random step:

```
// rand/walk_lib_pass.cpp
void random_step( vector<float>& step ) {
    for (;;) {
        random_coordinate(step);
    }
}
```

10.

This process of passing the arrays in stops at the *step* method, which we want to keep parameterless. So we add an option *cache* to the constructor to store the step vector as well as the position:

Code:

```
1 // rand/walk_lib_pass.cpp
2 class Mosquito {
3 private:
4     vector<float> pos;
5     vector<float> inc;
6     bool cache;
7 public:
8     Mosquito( int d,bool cache=false )
9         : pos( vector<float>(d,0.f)
10             ),cache(cache) {
11         if (cache) inc =
12             vector<float>(d,0.f);
13     };
14 }
```

Output:

```
D=3 after 10000 steps,
    distance= 76.7711
D=3 after 100000 steps,
    distance= 257.19
D=3 after 1000000
    steps, distance=
    956.122
run took: 2852
    milliseconds
D=3 after 10000 steps,
    distance= 87.034
D=3 after 100000 steps,
    distance= 256.655
D=3 after 1000000
    steps, distance=
    912.033
run took: 1762
```

11.

```
// rand/walk_lib_pass.cpp
void step() {
    int d = pos.size();
    if (cache) {
        random_step(inc);
        step( inc );
    } else {
        vector<float> incr(d);
        random_step(incr);
        step( incr );
    }
};
```

12. Sum of squares

There is still a problem with the *length* calculation. Since there is no reduction operator for 'sum of squares', we need to create a temporary vector for the squares,
(or do we?)
so that we can do a plus-reduction on it.

13. Exercise

Explore options for this temporary. Discuss what's most elegant, and measure performance improvement.

- This temporary can be passed in as a parameter;
- it can be stored in a global variable;
- or we can declare it `static`.
- With the C++20 standard, you could also use the `ranges` header.

14.

```
// rand/walk_lib_pass.cpp
float length( const vector<float>& step ) {
    vector<float> square;
    int s = step.size();
    if (square.size()!=s) square.resize(s);
    for ( int i=0; i<s; ++i) square[i] = step[i];
    for_each( square.begin(),square.end(),
              [] (float& x) { x *= x; } );
    auto l = sqrt
        ( accumulate( square.begin(),square.end(),0.f ) );
    return l;
};
```

15.

```
// rand/walk_lib_arr.cpp
template<int d>
float length( const array<float,d>& step ) {
    array<float,d> square = step;
    for_each( square.begin(),square.end(),
              [] (float& x) { x *= x; } );
    auto l = sqrt
        ( std::accumulate( square.begin(),square.end(),0.f ) );
    return l;
};
```


16. Optimization

While above we have removed all unnecessary allocation, we get an extra performance boost from optimizations from the compiler knowing the length of the array. Thus, instead of a loop of length two, the compiler will probably replace this by two explicit instructions.

17.

Code:

```
1 // rand/walk_arr.cpp
2 float avg_dist{0.f};
3 for ( int x=0; x<experiments; ++x )
4 {
5     Mosquito<dim> m;
6     for (int step=0; step<steps;
7         ++step)
8         m.step();
9     avg_dist += m.distance();
10 }
11 avg_dist /= experiments;
```

Output:

```
D=3 after 10000 steps,
    distance= 76.3221
D=3 after 100000 steps,
    distance= 247.5
D=3 after 1000000
    steps, distance=
    959.735
product took: 358
    milliseconds
```